

INFS2200 Report - 4582038

Task 1.

Q1.

Statement:

```
SELECT OWNER, CONSTRAINT_NAME, TABLE_NAME, SEARCH_CONDITION, INDEX_NAME
FROM USER_CONSTRAINTS
WHERE TABLE_NAME IN ('FILM_ACTOR', 'FILM', 'CATEGORY', 'LANGUAGE', 'FILM_CATEGORY', 'ACTOR');
```

Answer:

- The following constraints defined in Table 1 have been created by running prjScript.sql:
 - FK_FILMID1
 - PK_FILMID
 - PK_ACTORID

Q2.

Statements:

PK_CATEGORYID:

```
ALTER TABLE CATEGORY ADD CONSTRAINT PK_CATEGORYID PRIMARY KEY (category_id);
```

PK_LANGUAGEID:

```
ALTER TABLE LANGUAGE ADD CONSTRAINT PK_LANGUAGEID PRIMARY KEY (language_id);
```

UN_DESCRIPTION:

```
ALTER TABLE FILM ADD CONSTRAINT UN_DESCRIPTION UNIQUE (description);
```

CK_FNAME:

```
ALTER TABLE ACTOR MODIFY first_name VARCHAR2(45) NOT NULL;
```

CK_LNAME:

```
ALTER TABLE ACTOR MODIFY last_name VARCHAR2(45) NOT NULL;
```

CK_CATNAME:

```
ALTER TABLE CATEGORY MODIFY name VARCHAR2(25) NOT NULL;
```

CK_LANNAME:

```
ALTER TABLE LANGUAGE MODIFY name VARCHAR2(20) NOT NULL;
```

CK_TITLE:

```
ALTER TABLE FILM MODIFY title VARCHAR2(255) NOT NULL;
```

CK_RELEASEYR:

```
ALTER TABLE FILM ADD CONSTRAINT CK_RELEASEYR
CHECK (release_year <= 2020);
```

CK_RATING:

```
ALTER TABLE FILM ADD CONSTRAINT CK_RATING
CHECK (rating IN ('G', 'PG', 'PG-13', 'R', 'NC-17'));
```

CK_SPLFEATURES:

```
ALTER TABLE FILM ADD CONSTRAINT CK_SPLFEATURES
CHECK (special_features IN ('Trailers', 'Commentaries', 'Deleted Scenes', 'Behind the Scenes', NULL));
```

FK_LANGUAGEID:

```
ALTER TABLE FILM ADD CONSTRAINT FK_LANGUAGEID
FOREIGN KEY (language_id) REFERENCES LANGUAGE (language_id);
```

FK_ORLANGUAGEID:

```
ALTER TABLE FILM ADD CONSTRAINT FK_ORLANGUAGEID
FOREIGN KEY (original_language_id) REFERENCES LANGUAGE (language_id);
```

FK_ACTORID:

```
ALTER TABLE FILM_ACTOR ADD CONSTRAINT FK_ACTORID
FOREIGN KEY (actor_id) REFERENCES ACTOR (actor_id);
```

FK_CATEGORYID:

```
ALTER TABLE FILM_CATEGORY ADD CONSTRAINT FK_CATEGORYID
FOREIGN KEY (category_id) REFERENCES CATEGORY (category_id);
```

FK_FILMID2:

```
ALTER TABLE FILM_CATEGORY ADD CONSTRAINT FK_FILMID2
FOREIGN KEY (film_id) REFERENCES FILM (film_id);
```

Task 2.

Q1.

Statement:

```
CREATE SEQUENCE "FILM_ID_SEQ" MINVALUE 20010 MAXVALUE 999999999999
INCREMENT BY 10 START WITH 20010;
```

Q2.

Statement:

```
CREATE OR REPLACE TRIGGER "BI_FILM_ID"
  BEFORE INSERT ON "FILM"
  FOR EACH ROW
BEGIN
  SELECT "FILM_ID_SEQ".NEXTVAL INTO :NEW.film_id FROM DUAL;
END;
/
```

Q3.

Statement:

```
CREATE OR REPLACE TRIGGER "BI_FILM_DESP"
  BEFORE INSERT ON "FILM"
  FOR EACH ROW
  WHEN (NEW.rating IS NOT NULL AND NEW.language_id IS NOT NULL AND NEW.original_language_id IS NOT NULL)
DECLARE
  rating_count int;
  new_language varchar2(20);
  release_language varchar2(20);
BEGIN
  SELECT COUNT(*) INTO rating_count
  FROM FILM
  WHERE FILM.rating = :NEW.rating;

  SELECT name INTO release_language
  FROM LANGUAGE
  WHERE :NEW.original_language_id = LANGUAGE.language_id;

  SELECT name INTO new_language
  FROM LANGUAGE
  WHERE :NEW.language_id = LANGUAGE.language_id;

  SELECT (:NEW.description || :NEW.rating || '-' || TO_CHAR(rating_count) ||
  ': Originally in ' || release_language || '. Re-released in ' || new_language || '.') INTO :NEW.description
  FROM DUAL;
END;
/
```

Task 3.

Q1.

Statement:

```
SELECT title, length
FROM FILM, FILM_CATEGORY, CATEGORY
WHERE FILM.film_id = FILM_CATEGORY.film_id
  AND FILM_CATEGORY.category_id = CATEGORY.category_id
  AND CATEGORY.name = 'Action'
  AND length = (SELECT MIN(length)
                FROM FILM, FILM_CATEGORY, CATEGORY
                WHERE FILM.film_id = FILM_CATEGORY.film_id
                  AND FILM_CATEGORY.category_id = CATEGORY.category_id
                  AND CATEGORY.name = 'Action');
```

Output:

TITLE	LENGTH
STAR HEARTBREAKERS	00046
SPEAKEASY SILVERADO	00046
CLASH JAWS	00046
PIZZA BIRDCAGE	00046
FIRE AUTUMN	00046
METROPOLIS DREAM	00046
VANISHING AIRPORT	00046
PEACH PARADISE	00046
DOORS AIRPORT	00046
HOLLOW DYNAMITE	00046

Q2.

Statement:

```
CREATE VIEW MIN_ACTION_ACTORS AS
SELECT DISTINCT(A.actor_id), A.first_name, A.last_name
FROM ACTOR A, (SELECT actor_id
                FROM FILM_ACTOR, FILM, FILM_CATEGORY, CATEGORY
                WHERE FILM.film_id = FILM_CATEGORY.film_id
                  AND FILM.film_id = FILM_ACTOR.film_id
                  AND FILM_CATEGORY.category_id = CATEGORY.category_id
                  AND CATEGORY.name = 'Action'
                  AND length = (SELECT MIN(F.length)
                                FROM FILM F, FILM_CATEGORY, CATEGORY
                                WHERE FILM.film_id = FILM_CATEGORY.film_id
                                  AND FILM_CATEGORY.category_id = CATEGORY.category_id
                                  AND CATEGORY.name = 'Action')) B
WHERE A.actor_id = B.actor_id;
```

Q3.

Statement:

```
CREATE VIEW V_ACTION_ACTORS_2012 AS
SELECT DISTINCT(A.actor_id), A.first_name, A.last_name
FROM ACTOR A, (SELECT actor_id
FROM FILM_ACTOR, FILM, FILM_CATEGORY, CATEGORY
WHERE FILM.film_id = FILM_CATEGORY.film_id
AND FILM.film_id = FILM_ACTOR.film_id
AND FILM_CATEGORY.category_id = CATEGORY.category_id
AND CATEGORY.name = 'Action'
AND FILM.release_year = '2012') B
WHERE A.actor_id = B.actor_id;
```

Q4.

Statement:

```
CREATE MATERIALIZED VIEW MV_ACTION_ACTORS_2012
BUILD IMMEDIATE
AS
SELECT DISTINCT(A.actor_id), A.first_name, A.last_name
FROM ACTOR A, (SELECT actor_id
FROM FILM_ACTOR, FILM, FILM_CATEGORY, CATEGORY
WHERE FILM.film_id = FILM_CATEGORY.film_id
AND FILM.film_id = FILM_ACTOR.film_id
AND FILM_CATEGORY.category_id = CATEGORY.category_id
AND CATEGORY.name = 'Action'
AND FILM.release_year = '2012') B
WHERE A.actor_id = B.actor_id;
```

Q5.

Statements To Retrieve Information:

```
SET TIMING ON;

SELECT * FROM V_ACTION_ACTORS_2012;
SELECT * FROM MV_ACTION_ACTORS_2012;

SET TIMING OFF;

EXPLAIN PLAN FOR SELECT * FROM V_ACTION_ACTORS_2012;
SELECT PLAN_TABLE_OUTPUT FROM TABLE (DBMS_XPLAN.DISPLAY);

EXPLAIN PLAN FOR SELECT * FROM MV_ACTION_ACTORS_2012;
SELECT PLAN_TABLE_OUTPUT FROM TABLE (DBMS_XPLAN.DISPLAY);
```


Virtual View:

- Elapsed:

Elapsed: 00:00:00.20

- Query Execution Plan:

```
PLAN_TABLE_OUTPUT
-----
Plan hash value: 3856098404
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		528	32208	222 (2)	00:00:01
1	VIEW	V_ACTION_ACTORS_2012	528	32208	222 (2)	00:00:01
2	HASH UNIQUE		528	77616	222 (2)	00:00:01
3	NESTED LOOPS		528	77616	221 (2)	00:00:01
4	NESTED LOOPS		528	77616	221 (2)	00:00:01
* 5	HASH JOIN		528	45408	221 (2)	00:00:01
* 6	HASH JOIN SEMI		21	1260	151 (1)	00:00:01
7	MERGE JOIN CARTESIAN		333	11322	139 (0)	00:00:01
* 8	TABLE ACCESS FULL	CATEGORY	1	27	3 (0)	00:00:01
9	BUFFER SORT		333	2331	136 (0)	00:00:01
* 10	TABLE ACCESS FULL	FILM	333	2331	136 (0)	00:00:01
11	TABLE ACCESS FULL	FILM_CATEGORY	20000	507K	11 (0)	00:00:01
12	TABLE ACCESS FULL	FILM_ACTOR	129K	3299K	69 (2)	00:00:01
* 13	INDEX UNIQUE SCAN	PK_ACTORID	1		0 (0)	00:00:01
14	TABLE ACCESS BY INDEX ROWID	ACTOR	1	61	0 (0)	00:00:01

Materialised View:

- Elapsed Time:

Elapsed: 00:00:00.18

- Query Execution Plan:

```
PLAN_TABLE_OUTPUT
-----
Plan hash value: 1015139828
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		109	1744	3 (0)	00:00:01
1	MAT_VIEW ACCESS FULL	MV_ACTION_ACTORS_2012	109	1744	3 (0)	00:00:01

Answer:

- As shown by the two elapsed time recordings, the materialised view slightly reduced the required time to process the query.
- Additionally, as shown by its respective query execution plan (QEP), the query on the virtual view is translated to another query on base tables with numerous joins and tables which significantly increases CPU cost and query process time. In comparison, a single full table access can be performed on the materialised view (allowing the query to be performed directly on the view). This significantly reduced the number of necessary operations which in turn reduced CPU cost and runtime of the query, when compared to the virtual view.

Task 4.

Q1.

Statement:

```
SELECT *
FROM FILM F
WHERE INSTR(F.description, 'Boat') IS NOT NULL AND INSTR(F.description, 'Boat') > 0
ORDER BY F.title
FETCH FIRST 100 ROWS ONLY;
```

Output (top 3 and bottom 3 tuples):

FILM_ID	TITLE	DESCRIPTION	RELEASE_YEAR	LANGUAGE_ID	ORIGINAL_LANGUAGE_ID	RENTAL_DURATION	RENTAL_RATE	LENGTH	REPLACEMENT_COST	RATING
SPECIAL_FEATURES										
05880	ACADEMY REDS	A Beautiful Story of a Crocodile and a Lumberjack who must Vanquish a Forensic Psychologist in a Jet Boat	1974	001		007	0000.99	00142	00017.99	R
Commentaries										
04952	ADAPTATION JEOPARDY	A Thrilling Saga of a Student and a Car who must Battle a Technical Writer in A U-Boat	1958	006		004	0002.99	00129	00028.99	R
Commentaries										
06107	AFFAIR PANIC	An Intrepid Yarn of a Teacher and a Student who must Pursue a Madman in a Jet Boat	1975	001		004	0000.99	00136	00023.99	NC-17
Commentaries										
09446	BEETHOVEN CIRCUS	An Intrepid Story of a Hunter and a Database Administrator who must Face a Monkey in a Jet Boat	1975	001		007	0004.99	00057	00027.99	PG
Commentaries										
SPECIAL_FEATURES										
18513	BEETHOVEN MOCKINGBIRD	A Taut Story of a Woman and a Man who must Find a Database Administrator in A U-Boat	1978	004		004	0000.99	00105	00016.99	NC-17
Trailers										
15677	BEHAVIOR SLIPPER	An Amazing Display of a Hunter and a Waitress who must Redeem a Pioneer in a Jet Boat	1955	001		006	0004.99	00084	00024.99	PG
Deleted Scenes										
100 rows selected.										

Q2.

Statement:

```
CREATE INDEX IDX_BOAT ON FILM(INSTR(description, 'Boat'));
```

Justification:

- As the query searches for tuples which satisfy a single condition (description contains 'Boat') a single column index using the function INSTR(description, 'Boat') indicates if the description contains 'Boat'. This simplifies the where condition of the query to check the index rather than performing INSTR for each tuple in FILM each time the query is queried.

Q3.

Statements To Retrieve Information:

```
DROP INDEX IDX_BOAT;
SET TIMING ON;

SELECT *
FROM FILM F
WHERE INSTR(F.description, 'Boat') IS NOT NULL AND INSTR(F.description, 'Boat') > 0
ORDER BY F.title
FETCH FIRST 100 ROWS ONLY;

SET TIMING OFF;

EXPLAIN PLAN FOR SELECT *
FROM FILM F
WHERE INSTR(F.description, 'Boat') IS NOT NULL AND INSTR(F.description, 'Boat') > 0
ORDER BY F.title
FETCH FIRST 100 ROWS ONLY;
SELECT PLAN_TABLE_OUTPUT FROM TABLE (DBMS_XPLAN.DISPLAY);

CREATE INDEX IDX_BOAT ON FILM(INSTR(description, 'Boat'));

SET TIMING ON;

SELECT *
FROM FILM F
WHERE INSTR(F.description, 'Boat') IS NOT NULL AND INSTR(F.description, 'Boat') > 0
ORDER BY F.title
FETCH FIRST 100 ROWS ONLY;

SET TIMING OFF;

EXPLAIN PLAN FOR SELECT *
FROM FILM F
WHERE INSTR(F.description, 'Boat') IS NOT NULL AND INSTR(F.description, 'Boat') > 0
ORDER BY F.title
FETCH FIRST 100 ROWS ONLY;
SELECT PLAN_TABLE_OUTPUT FROM TABLE (DBMS_XPLAN.DISPLAY);
```


Before Index:

- Elapsed Time:

Elapsed: 00:00:00.50

- Execution Plan:

```
PLAN_TABLE_OUTPUT
-----
Plan hash value: 2865634321

-----
| Id | Operation                      | Name | Rows  | Bytes | Cost (%CPU)| Time     |
-----
|  0 | SELECT STATEMENT                |      |    100 | 63900 |    138 (2) | 00:00:01 |
| * 1 | VIEW                            |      |    100 | 63900 |    138 (2) | 00:00:01 |
| * 2 | WINDOW SORT PUSHED RANK        |      |    1000 | 142K |    138 (2) | 00:00:01 |
| * 3 | TABLE ACCESS FULL             | FILM |    1000 | 142K |    137 (1) | 00:00:01 |
-----

Predicate Information (identified by operation id):
-----
 1 - filter("from$_subquery$_002"."rowlimit_$$_rownumber"<=100)
 2 - filter(ROW_NUMBER() OVER ( ORDER BY "F"."TITLE")<=100)
 3 - filter(INSTR("F"."DESCRIPTION",'Boat')>0 AND
        INSTR("F"."DESCRIPTION",'Boat') IS NOT NULL)
```

After Index:

- Elapsed Time:

Elapsed: 00:00:00.45

- Execution Plan:

```
PLAN_TABLE_OUTPUT
-----
Plan hash value: 2388608894

-----
| Id | Operation                      | Name      | Rows  | Bytes | Cost (%CPU)| Time     |
-----
|  0 | SELECT STATEMENT                |           |    100 | 63900 |    22 (5) | 00:00:01 |
| * 1 | VIEW                            |           |    100 | 63900 |    22 (5) | 00:00:01 |
| * 2 | WINDOW SORT PUSHED RANK        |           |    1000 | 151K |    22 (5) | 00:00:01 |
|  3 | TABLE ACCESS BY INDEX ROWID BATCHED | FILM      |    1000 | 151K |    21 (0) | 00:00:01 |
| * 4 | INDEX RANGE SCAN                | IDX_BOAT  |     180 |      |     2 (0) | 00:00:01 |
-----

Predicate Information (identified by operation id):
-----
 1 - filter("from$_subquery$_002"."rowlimit_$$_rownumber"<=100)
 2 - filter(ROW_NUMBER() OVER ( ORDER BY "F"."TITLE")<=100)
 4 - access(INSTR("DESCRIPTION",'Boat')>0)
```

Answer:

- As shown by the above elapsed times, the index speeds up query processing.
- Additionally, as seen by the corresponding execution plans, the index significantly reduced CPU cost. This can be attributed to the fact that the indexed table can be accessed by index instead of a full table access (accessing only tuples which meets the condition).
- Thus, both the recorded elapsed times and execution plans indicate a notable decrease in query processing time.

Q4.

Statement:

```
SELECT COUNT(*)
FROM (SELECT film_id
      FROM FILM G, (SELECT F.release_year, F.rating, F.special_features, COUNT(*)
                    FROM FILM F
                    GROUP BY F.release_year, F.rating, F.special_features
                    HAVING COUNT(*) >= 41) H
      WHERE G.release_year = H.release_year AND
            G.rating = H.rating AND
            G.special_features = H.special_features);
```

Output:

```
COUNT(*)
-----
      265
```

Q5.

Answer:

- The most suitable index type to create on the columns is a composite secondary index.
- A composite index should be used because multiple columns are used within the group by clause, allowing for easier grouping of equivalent element tuples.
- A secondary index should be used since the table is ordered on the primary key and as such another primary/clustering index cannot be created.
- Additionally, another possible option would be a bitmap index, however due to the possibly large number of unique values it is not applicable in this situation.
- A function-based index was considered but was rejected as it would significantly affect insertion, update and deletion procedures as all entries in the index would need to be updated.
- As such, a composite secondary index seems the most applicable in this situation.

Task 5.

Q1.

Statements:

```
ANALYZE INDEX PK_FILMID VALIDATE STRUCTURE;
```

---- a, b)

```
SELECT HEIGHT, LF_BLKs
FROM INDEX_STATS;
```

---- c)

```
SELECT BLOCKS
FROM USER_TABLES
WHERE table_name = 'FILM';
```

Outputs:

a, b)

HEIGHT	LF_BLKs
2	37

c)

BLOCKS
496

Answers:

- a) Height of B+ Tree = 2
- b) Number of leaf blocks in the B+ tree index = 37
- c) Number of block access for a full table scan on 'film' = 496

Q2.

Statement To Retrieve QEP:

```
EXPLAIN PLAN FOR SELECT /*+RULE*/ * FROM FILM WHERE FILM_ID > 100;
```

Output:

```
PLAN_TABLE_OUTPUT
-----
Plan hash value: 657888726

-----
| Id | Operation                      | Name |
-----
| 0  | SELECT STATEMENT                |      |
| 1  | TABLE ACCESS BY INDEX ROWID    | FILM |
|* 2  | INDEX RANGE SCAN                | PK_FILMID |
-----

Predicate Information (identified by operation id):
-----
   2 - access("FILM_ID">19990)

Note
-----
   - rule based optimizer used (consider using cbo)

18 rows selected.
```

Answer:

- As shown by the above QEP, the query first performs an index range scan on the primary key of FILM for film_ids > 100. Then locate the related tuples in FILM using the indices. Then finally performs the select statement and returns the requested fields from the resultant tuples.

Q3.

Statement To Retrieve QEP:

```
EXPLAIN PLAN FOR SELECT * FROM FILM WHERE FILM_ID > 100;
```

Output:

```
PLAN_TABLE_OUTPUT
-----
Plan hash value: 1232367652

-----
| Id | Operation          | Name | Rows  | Bytes | Cost (%CPU)| Time     |
-----|-----|-----|-----|-----|-----|-----|
|  0 | SELECT STATEMENT   |      |  19901 | 2837K |    137   (1)| 00:00:01 |
|*  1 | TABLE ACCESS FULL| FILM |  19901 | 2837K |    137   (1)| 00:00:01 |
-----

Predicate Information (identified by operation id):
-----
   1 - filter("FILM_ID">100)

13 rows selected.
```

Answer:

- The query first performs a full table access on FILM checking all tuples against the arguments in the where clause. Then performs the select statement returning the necessary fields of the resultant tuples.
- An obvious difference between the plans obtained in Q3 and Q2 is that Q2 utilises the primary index to search for relevant tuples while Q3 just performs a full table search checking each tuple.
- In Q2, the index range scan performs tree height – 1 + 2 * leaf blocks (at most) block accesses (in this case 2*37 accesses). The true number of block accesses index range scan would perform is tree height – 1 + # of matching leaf blocks + # of matching data blocks. On the other hand in Q3, the full table scan performs exactly 496 block accesses. The main difference between the two plans is the resultant average number of block accesses.
- Another significant difference, is the number of expected rows which must be accessed and returned.
- Note:
 - o For Q2, the index range scan might the same block multiple times if the number of valid row ids is significant, which would increase the number of block access performed.

Q4.

Statement To Retrieve QEP:

```
EXPLAIN PLAN FOR SELECT * FROM FILM WHERE FILM_ID > 19990;
```

Output:

```
PLAN_TABLE_OUTPUT
-----
Plan hash value: 1620599584

-----
| Id | Operation                      | Name      | Rows  | Bytes | Cost (%CPU)| Time     |
-----
|  0 | SELECT STATEMENT                |           |      10 | 1460 |      3 (0)| 00:00:01 |
|  1 | TABLE ACCESS BY INDEX ROWID BATCHED | FILM      |      10 | 1460 |      3 (0)| 00:00:01 |
|*  2 | INDEX RANGE SCAN                 | PK_FILMID |      10 |      |      2 (0)| 00:00:01 |
-----

Predicate Information (identified by operation id):
-----
   2 - access("FILM_ID">19990)

14 rows selected.
```

Answer:

- The query first performs an index range scan on PK_FILM for values > 19990. It uses the indices to performing multi-block read on a few tuples which satisfy the where clause, then tries to access rows in block order. Then performs the select statement returning the necessary fields of the resultant tuples.
- The main observable difference between the plans produced by Q3 and Q4 is that Q4 utilises the primary index whilst Q3 does not.
- The number of block accesses performed by index range scan, in Q4, is tree height – 1 + # of matching leaf blocks + # of matching data blocks. While Q3 performs exactly 496 block accesses. Thus Q4's QEP has a significantly reduced cost compared to Q3's QEP.
- Q4 utilises the index because it is trying to access a significantly smaller subset of tuples than those accessed in Q3.

Q5.

Statement To Retrieve QEP:

```
EXPLAIN PLAN FOR SELECT * FROM FILM WHERE FILM ID = 100;
```

Output:

```
PLAN_TABLE_OUTPUT
-----
Plan hash value: 2104374699

-----
| Id | Operation                      | Name      | Rows  | Bytes | Cost (%CPU)| Time     |
-----
|  0 | SELECT STATEMENT                |           |      1 | 146 |      2 (0)| 00:00:01 |
|  1 | TABLE ACCESS BY INDEX ROWID    | FILM      |      1 | 146 |      2 (0)| 00:00:01 |
|*  2 | INDEX UNIQUE SCAN                | PK_FILMID |      1 |      |      1 (0)| 00:00:01 |
-----

Predicate Information (identified by operation id):
-----
   2 - access("FILM_ID"=100)

14 rows selected.
```


Answer:

- The query first uses PK_FILMID in a unique scan operation to evaluate the where clause criteria, returning exactly one row-id from the index.
- Rows are then accessed using the index.
- Then the select statement returns the resultant rows.
- Differences between the QEP of Q5 and Q3:
 - Q5 performs tree height + 1 block access and only returns one row (since FILM_ID is the primary key) unlike Q3 which accesses every row and block. Similar to Q4, Q5 has a significantly lower cost than Q3.
 - Similarly to Q4, Q5 utilises the index because it is trying to access a significantly smaller subset of tuples than those accessed in Q3.